

```
1 from sklearn.metrics import accuracy_score
2 import matplotlib.pyplot as plt
3 import numpy as np
4 import struct
5 import pickle
6 import time
7 import os
```

```
1 root = ''
2 data_path = os.path.join(root, 'data')
```

▼ Layer Object

```
1 class Layer:
2     def __init__(self, layer_in_size: int, layer_out_size: int, activation: str):
3         np.random.seed(42)
4
5         # Drawing samples from LeCun normal distribution
6         # Source: https://arxiv.org/pdf/1706.02515.pdf
7
8         self.weights = np.random.normal(loc=0, scale=(1 / layer_in_size), size=(layer_in_size, layer_out_size))
9         self.biases = np.random.normal(loc=0, scale=(1 / layer_in_size), size=(layer_out_size, 1))
10
11         if activation == 'tanh':
12             self.__activation_function = lambda v: np.tanh(v)
13             self.__derivative_function = lambda v: 1 - (np.tanh(v) ** 2)
14         elif activation == 'sigmoid':
15             self.__activation_function = lambda v: 1 / (1 + np.exp(-v))
16             self.__derivative_function = lambda v: self.__sigmoid_derivative(v)
17         else:
18             self.__activation_function = lambda v: v
19             self.__derivative_function = lambda v: 1
20
21     def local_fields(self, data_in):
22         return self.weights.T.dot(data_in) + self.biases
23
24     def activations(self, local_fields):
25         return self.__activation_function(local_fields)
26
27     def derivatives(self, local_fields):
28         return self.__derivative_function(local_fields)
29
30     @staticmethod
31     def __sigmoid_derivative(v):
32         a = 1 / (1 + np.exp(-v))
33         return a * (1 - a)
```

▼ Simple Neural Network

```
1 class NeuralNetwork:
2     def __init__(self,
3                 data_x: np.ndarray,
4                 data_y: np.ndarray,
5                 test_x: np.ndarray,
6                 test_y: np.ndarray,
7                 hidden_layers: tuple,
8                 learning_rate: float,
9                 lr_decay_factor: float = 0.9):
10         self.data = data_x
11         self.labels = data_y
12
13         self.test_x = test_x
14         self.test_y = test_y
15
16         self.learning_rate = learning_rate
17         self.decay_factor = lr_decay_factor
18
19         self.n_features = self.data.shape[0]
20         self.n_outputs = self.labels.shape[0]
21
22         self.n_samples = self.data.shape[1]
23
24         # Layers of neural network
```

```

25     self.nn_layers = list()
26     self.nn_layers.append(Layer(self.n_features, hidden_layers[0]['num_nodes'], hidden_layers[0]['activation']))
27     for i in range(1, len(hidden_layers)):
28         self.nn_layers.append(Layer(hidden_layers[i - 1]['num_nodes'], hidden_layers[i]['num_nodes'],
29                                     hidden_layers[i]['activation']))
30     self.nn_layers.append(Layer(hidden_layers[-1]['num_nodes'], self.n_outputs, 'sigmoid'))
31
32     def calc_stats(self, data, labels):
33         predictions = self.predict(data)
34         mse = np.sum(((predictions - labels) ** 2) / labels.shape[1])
35         acc = accuracy_score(np.argmax(labels, axis=0), np.argmax(predictions, axis=0))
36         return mse, acc
37
38     def predict(self, data):
39         local_fields, activations = self.__forward(data)
40         return activations[-1]
41
42     def __forward(self, x_i):
43         local_fields = list()
44         activations = list()
45         current_input = x_i
46         for layer in self.nn_layers:
47             z = layer.local_fields(current_input)
48             local_fields.append(z)
49             a = layer.activations(z)
50             activations.append(a)
51             current_input = a
52         return local_fields, activations
53
54     def __backward(self, initial_delta, local_fields):
55         current_delta = initial_delta
56         layer_delta = list()
57         for i in reversed(range(len(self.nn_layers))):
58             if i == len(self.nn_layers) - 1:
59                 delta = current_delta * self.nn_layers[i].derivatives(local_fields[i])
60             else:
61                 delta = self.nn_layers[i + 1].weights.dot(current_delta) * self.nn_layers[i].derivatives(
62                     local_fields[i])
63             layer_delta.insert(0, delta)
64             current_delta = delta
65         return layer_delta
66
67     def __update_layer_params(self, x_i, layer_delta, activations):
68         current_input = x_i
69         for layer, activation, delta in zip(self.nn_layers, activations, layer_delta):
70             layer.weights = layer.weights - self.learning_rate * current_input.dot(delta.T)
71             layer.biases = layer.biases - self.learning_rate * delta
72             current_input = activation
73
74     def train(self):
75         train_epoch_stats = list()
76         test_epoch_stats = list()
77         train_mse, train_acc = self.calc_stats(self.data, self.labels)
78         train_epoch_stats.append([0, train_mse, train_acc])
79         test_mse, test_acc = self.calc_stats(self.test_x, self.test_y)
80         test_epoch_stats.append([0, test_mse, test_acc])
81         epoch_cnt = 1
82         while test_epoch_stats[-1][2] < 0.955:
83             start_time = time.time()
84             for i in range(self.n_samples):
85                 x_i = self.data[:, i].reshape((self.n_features, 1))
86                 d_i = self.labels[:, i].reshape((self.n_outputs, 1))
87                 local_fields, activations = self.__forward(x_i)
88                 y_i = activations[-1]
89                 initial_delta = 2 * (y_i - d_i) / self.n_samples
90                 delta_list = self.__backward(initial_delta, local_fields)
91                 self.__update_layer_params(x_i, delta_list, activations)
92                 train_mse, train_acc = self.calc_stats(self.data, self.labels)
93                 test_mse, test_acc = self.calc_stats(self.test_x, self.test_y)
94                 print(
95                     '[Epoch: {}] => Train MSE: {:.4f}, Train Accuracy: {:.4f}, Test Accuracy: {:.4f}, Epoch Duration: {:.4f} S'.format
96                     epoch_cnt, train_mse, train_acc, test_acc, time.time() - start_time))
97
98             if test_acc <= test_epoch_stats[-1][2]:
99                 self.learning_rate = self.learning_rate * self.decay_factor
100
101         train_epoch_stats.append([epoch_cnt, train_mse, train_acc])
102         test_epoch_stats.append([epoch_cnt, test_mse, test_acc])
103         epoch_cnt += 1
104         return np.array(train_epoch_stats), np.array(test_epoch_stats)

```

```

105
106     def save_params(self, save_path):
107         nn_params = list()
108         for layer in self.nn_layers:
109             nn_params.append({
110                 'weight': layer.weights,
111                 'bias': layer.biases
112             })
113         with open(save_path, 'wb') as model_params:
114             pickle.dump(nn_params, model_params)
115
116     def load_params(self, params_path):
117         nn_params = pickle.load(open(params_path, 'rb'))
118         assert len(self.nn_layers) == len(nn_params)
119         for layer, params in zip(self.nn_layers, nn_params):
120             layer.weights = params['weight']
121             layer.biases = params['bias']

```

▼ Reading Data

```

1 def read_idx(filename):
2     with open(filename, 'rb') as f:
3         zero, data_type, dims = struct.unpack('>HBB', f.read(4))
4         shape = tuple(struct.unpack('>I', f.read(4))[0] for d in range(dims))
5         return np.frombuffer(f.read(), dtype=np.uint8).reshape(shape)

```

```

1 train_images = read_idx(os.path.join(data_path, 'train-images-idx3-ubyte'))
2 train_images = train_images.reshape((train_images.shape[0], train_images.shape[1] * train_images.shape[2])).T
3
4 train_labels = read_idx(os.path.join(data_path, 'train-labels-idx1-ubyte'))
5 train_labels = np.eye(10)[train_labels].T
6
7 test_images = read_idx(os.path.join(data_path, 't10k-images-idx3-ubyte'))
8 test_images = test_images.reshape((test_images.shape[0], test_images.shape[1] * test_images.shape[2])).T
9
10 test_labels = read_idx(os.path.join(data_path, 't10k-labels-idx1-ubyte'))
11 test_labels_norm = test_labels.copy()
12 test_labels = np.eye(10)[test_labels].T

```

▼ Training Process

```

1 hidden_layer_params = (
2     {'num_nodes': 256,
3     'activation': 'tanh'},
4     {'num_nodes': 16,
5     'activation': 'tanh'}
6 )
7
8 nn_regressor = NeuralNetwork(train_images,
9                             train_labels,
10                             test_images,
11                             test_labels,
12                             hidden_layers=hidden_layer_params,
13                             learning_rate=12,
14                             lr_decay_factor=0.7)
15 train_epoch_stats, test_epoch_stats = nn_regressor.train()
16 nn_regressor.save_params(os.path.join(root, 'model_params/256_16_sigmoid_params.pkl'))
17
18 # nn_regressor.load_params(os.path.join(root, 'model_params/256_16_sigmoid_params.pkl'))
19 # with open(os.path.join(root, 'model_stats/256_16_sigmoid_stats.pkl'), 'wb') as stats_file:
20 #     pickle.dump((train_epoch_stats, test_epoch_stats), stats_file)

```



[Epoch: 1] => Train MSE: 0.8954, Train Accuracy: 0.2083, Test Accuracy: 0.2112, Epoch Duration: 70.6459 S
[Epoch: 2] => Train MSE: 0.8654, Train Accuracy: 0.2059, Test Accuracy: 0.2081, Epoch Duration: 71.3460 S
[Epoch: 3] => Train MSE: 0.8501, Train Accuracy: 0.2122, Test Accuracy: 0.2106, Epoch Duration: 69.2662 S
[Epoch: 4] => Train MSE: 0.8395, Train Accuracy: 0.2127, Test Accuracy: 0.2110, Epoch Duration: 73.9732 S
[Epoch: 5] => Train MSE: 0.8259, Train Accuracy: 0.2136, Test Accuracy: 0.2116, Epoch Duration: 70.0234 S
[Epoch: 6] => Train MSE: 0.7818, Train Accuracy: 0.3125, Test Accuracy: 0.3109, Epoch Duration: 69.6082 S
[Epoch: 7] => Train MSE: 0.7173, Train Accuracy: 0.5504, Test Accuracy: 0.5507, Epoch Duration: 70.4475 S
[Epoch: 8] => Train MSE: 0.6012, Train Accuracy: 0.6069, Test Accuracy: 0.6019, Epoch Duration: 68.0290 S
[Epoch: 9] => Train MSE: 0.4829, Train Accuracy: 0.7888, Test Accuracy: 0.7924, Epoch Duration: 72.5066 S
[Epoch: 10] => Train MSE: 0.3897, Train Accuracy: 0.8565, Test Accuracy: 0.8570, Epoch Duration: 68.1072 S
[Epoch: 11] => Train MSE: 0.3233, Train Accuracy: 0.8991, Test Accuracy: 0.8964, Epoch Duration: 72.3832 S
[Epoch: 12] => Train MSE: 0.2719, Train Accuracy: 0.9138, Test Accuracy: 0.9123, Epoch Duration: 68.2074 S
[Epoch: 13] => Train MSE: 0.2379, Train Accuracy: 0.9222, Test Accuracy: 0.9162, Epoch Duration: 66.5742 S
[Epoch: 14] => Train MSE: 0.2119, Train Accuracy: 0.9279, Test Accuracy: 0.9279, Epoch Duration: 69.1238 S
[Epoch: 15] => Train MSE: 0.1976, Train Accuracy: 0.9279, Test Accuracy: 0.9273, Epoch Duration: 66.5518 S
[Epoch: 16] => Train MSE: 0.1723, Train Accuracy: 0.9391, Test Accuracy: 0.9376, Epoch Duration: 70.3178 S
[Epoch: 17] => Train MSE: 0.1588, Train Accuracy: 0.9438, Test Accuracy: 0.9384, Epoch Duration: 66.7519 S
[Epoch: 18] => Train MSE: 0.1509, Train Accuracy: 0.9457, Test Accuracy: 0.9447, Epoch Duration: 70.9285 S
[Epoch: 19] => Train MSE: 0.1420, Train Accuracy: 0.9479, Test Accuracy: 0.9426, Epoch Duration: 67.6233 S
[Epoch: 20] => Train MSE: 0.1298, Train Accuracy: 0.9539, Test Accuracy: 0.9450, Epoch Duration: 65.8898 S
[Epoch: 21] => Train MSE: 0.1270, Train Accuracy: 0.9543, Test Accuracy: 0.9485, Epoch Duration: 67.8939 S
[Epoch: 22] => Train MSE: 0.1187, Train Accuracy: 0.9575, Test Accuracy: 0.9503, Epoch Duration: 65.7243 S
[Epoch: 23] => Train MSE: 0.1173, Train Accuracy: 0.9576, Test Accuracy: 0.9508, Epoch Duration: 68.5207 S
[Epoch: 24] => Train MSE: 0.1121, Train Accuracy: 0.9602, Test Accuracy: 0.9530, Epoch Duration: 65.5310 S
[Epoch: 25] => Train MSE: 0.1091, Train Accuracy: 0.9597, Test Accuracy: 0.9522, Epoch Duration: 70.2115 S
[Epoch: 26] => Train MSE: 0.1033, Train Accuracy: 0.9631, Test Accuracy: 0.9528, Epoch Duration: 65.8000 S
[Epoch: 27] => Train MSE: 0.0988, Train Accuracy: 0.9650, Test Accuracy: 0.9548, Epoch Duration: 65.7438 S
[Epoch: 28] => Train MSE: 0.0955, Train Accuracy: 0.9666, Test Accuracy: 0.9559, Epoch Duration: 66.9296 S

▼ Predicting random samples from test data just to make sure

```
1 random_indices = np.random.randint(0, 10000, size=(100,))
2 test_preds = nn_regressor.predict(test_images[:, random_indices])
3 test_preds = np.argmax(test_preds, axis=0)
4 print('Test accuracy: {:.4f}'.format(accuracy_score(test_labels_norm[random_indices], test_preds)))
```

📄 Test accuracy: 0.9600

▼ Network Architecture

The network contains two hidden layers with 256 and 16 neurons respectively. Activation function for hidden layers is 'tanh' while the output layer uses a 'sigmoid' activation.

▼ Output Representation

The output is represented as one-hot encoded vectors corresponding to a single digit.
For example, the digit 3 is represented as [0, 0, 0, 1, 0, 0, 0, 0, 0, 0] and 9 is represented as [0, 0, 0, 0, 0, 0, 0, 0, 1]

▼ Activation and Learning Rate

Hidden layer neurons use 'tanh' activation while output layer uses 'sigmoid' activation.
The initial learning rate is 12 for all neurons and it is decayed by a factor of 0.7 when there is a decrease in test accuracy from the previous epoch.

▼ Energy Function

The energy function used is the Mean Squared Error because of its convexity. As a result, there is only one global optimum, a smooth loss landscape and gradients can be found easily.

$$\frac{1}{n} \sum_{i=1}^n (d_i - y_i)^2$$

▼ Hyperparameters

Hyperparameter	Value	Reason
Learning rate	12	Moderately large learning rates can help faster convergence in the initial training steps. Such learning rates are usually used in combination with a decay rate less than 1
Decay rate	0.7	Learning rate is decayed at a rate of 0.7 so as to avoid divergence

Number of hidden layers	2	Having more hidden layer increases neural network complexity allowing it to represent more complex functions. Our application requires only moderate complexity and thus 2 hidden layers. This also helps avoid drastic reduction in the size of inputs to subsequent layers.
Number of Hidden layer Neurons	Powers of 2 256 and 16	Having powers-of-two neurons is a general recommendation. Some sources state this helps easier compiler optimizations.
Activations	tanh, tanh and sigmoid	In Efficient BackProp the authors argue that convergence is faster if the average of inputs to each layer is close to zero. Hyperbolic tangent function has this property where the outputs $\in (-1, 1)$ thus pushing the average close to zero. The output activation is sigmoid because it more closely represents probabilities. However, tanh also gives good performance.

▼ Design Details

Hidden Layers - Initially the network contained one hidden layer with 64 neurons and tanh activations for all layers. However, this network plateaued at 90% training accuracy indicating a more complex model could perform better.

Learning Rates - A couple of learning rates have also been tried allowing me to conclude lower learning rates caused the initial convergence to be slow thus taking large number of epochs to converge. Learning rates greater than 20 caused divergence

Decay Rates - Higher decay rates cause the network to oscillate around a single point (error) while lower decay rates resulted in slow convergence.

Activation Functions - Tanh activation for output layer resulted in better performance and higher convergence rate. However, sigmoid activation allowed the network to train in a more stable manner. Also, sigmoid is a better representation for probabilities than a tanh.


Final Design

Two hidden layers with 256 and 16 neurons respectively and tanh activation. Output layer with sigmoid activation. Initial learning rate of 12 and a decay factor of 0.7.

▼ Plots

▼ MSE Plots

```
1 fig = plt.figure(figsize=(15, 4))
2 plt.subplot(1, 3, 1)
3 plt.plot(train_epoch_stats[:, 0], train_epoch_stats[:, 1])
4 plt.xlabel('Epoch')
5 plt.ylabel('Train MSE')
6 plt.subplot(1, 3, 2)
7 plt.plot(test_epoch_stats[:, 0], test_epoch_stats[:, 1])
8 plt.xlabel('Epoch')
9 plt.ylabel('Test MSE')
10 plt.subplot(1, 3, 3)
11 plt.plot(train_epoch_stats[:, 0], train_epoch_stats[:, 1], label='Train MSE')
12 plt.plot(test_epoch_stats[:, 0], test_epoch_stats[:, 1], label='Test MSE')
13 plt.legend(loc='best')
14 plt.xlabel('Epoch')
15 plt.ylabel('Test MSE')
```



▼ Error Plots


```
1 fig = plt.figure(figsize=(15, 4))
2 plt.subplot(1, 3, 1)
3 plt.plot(train_epoch_stats[:, 0], 1-train_epoch_stats[:, 2])
4 plt.xlabel('Epoch')
5 plt.ylabel('Train Error')
6 plt.title('Training Error')
7 plt.subplot(1, 3, 2)
```

```

8 plt.plot(test_epoch_stats[:, 0], 1-test_epoch_stats[:, 2])
9 plt.xlabel('Epoch')
10 plt.ylabel('Test Error')
11 plt.title('Test Error')
12 plt.subplot(1, 3, 3)
13 plt.plot(train_epoch_stats[:, 0], 1-train_epoch_stats[:, 2], label='Train Error')
14 plt.plot(test_epoch_stats[:, 0], 1-test_epoch_stats[:, 2], label='Test Error')
15 plt.legend(loc='best')
16 plt.xlabel('Epoch')
17 plt.ylabel('Error')
18 plt.title('Error')

```

➤ Text(0.5, 1.0, 'Error')




▼ Accuracy Plots

```

1 fig = plt.figure(figsize=(15, 4))
2 plt.subplot(1, 3, 1)
3 plt.plot(train_epoch_stats[:, 0], train_epoch_stats[:, 2])
4 plt.xlabel('Epoch')
5 plt.ylabel('Train Accuracy')
6 plt.title('Train Accuracy')
7 plt.subplot(1, 3, 2)
8 plt.plot(test_epoch_stats[:, 0], test_epoch_stats[:, 2])
9 plt.xlabel('Epoch')
10 plt.ylabel('Test Accuracy')
11 plt.title('Test Accuracy')
12 plt.subplot(1, 3, 3)
13 plt.plot(train_epoch_stats[:, 0], train_epoch_stats[:, 2], label='Train Accuracy')
14 plt.plot(test_epoch_stats[:, 0], test_epoch_stats[:, 2], label='Test Accuracy')
15 plt.legend(loc='best')
16 plt.xlabel('Epoch')
17 plt.ylabel('Accuracy')
18 plt.title('Accuracy')

```

➤ Text(0.5, 1.0, 'Accuracy')



▼ Pseudocode

while test_accuracy < 0.955:

 for each training_example:

 for each layer i :

$v^i = weights^{iT} \cdot activation^{i-1}$ { $activation^{i-1}$ for first layer is the train example}

$a^i = f^i(v^i)$

(store v and a in an array)

for each layer i :

$$\delta^i = (weights^{i+1} \cdot \delta^{i+1}) * f^{i'}(v^i)$$

$$weights^i \leftarrow weights^i - \eta * (activation^{i-1} \cdot \delta^i)$$

$$bias^i \leftarrow bias^i - \eta * \delta^i$$

$$mse = \frac{1}{n} \sum (d_i - y_i)^2$$

$$test_accuracy = \frac{num_correct_test}{num_samples_test}$$

if test_accuracy < test_accuracy_prev:

$$\eta \leftarrow \eta * decay_factor$$

Where,

$v^i \rightarrow$ Local field values of i^{th} layer

$w^i, b^i \rightarrow$ Weights and bias respectively for layer i

$\delta^i \rightarrow$ Delta or error w.r.t layer i (Derivative of cost function w.r.t to layer i)

$f^i \rightarrow$ Activation function of layer i

$f^{i'} \rightarrow$ Derivarive of activation function of layer i

$a^i \rightarrow$ Activation of layer i

$\eta \rightarrow$ Learning rate

$\cdot \rightarrow$ Dot product

$*$ \rightarrow Element-wise multiplication